

Deep learning and applications – Part 3

Sylvain Lobry

13/10/2023

Before we start...

- <https://app.wooclap.com/VMIAI3>

Today's menu

- 1 Optimisation
- 2 Sequence modeling
- 3 Visual Question Answering
- 4 Activity

Before we start...

- You should be familiar with the concepts of:
 - Convolutional layers
 - Pooling layers (max and average)
 - Activation function
 - Loss function
 - Back propagation
 - Gradient descent

Optimization and tricks

Stochastic Gradient Descent

Reminder: computation of the gradient (in the supervised learning case):

We have l annotated samples $\{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_l, y_l)\}$ and the empirical risk is defined as:

$$J(f) = \frac{1}{l} \sum_{i=1}^l L(y_i, f(\mathbf{x}_i))$$

Where $L(y, \hat{y})$ is the loss function for one sample.

To compute the gradient, we would have:

$$\nabla J(f) = \frac{1}{l} \sum_{i=1}^l \nabla L(y_i, f(\mathbf{x}_i))$$

We can see that this is in $\theta(l)$: when the dataset grows, the computation of the gradient grows linearly.

Solution: **sample** $(\mathbf{x}_i, y_i) \in \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_l, y_l)\}$ and do a gradient descent step based on this.

Note: this solution is unbiased (the expectation is the same)

Optimization and tricks

Stochastic Gradient Descent

However, in the case of deep learning: billions of parameters to update.

	Gradient descent	Batch SGD	Stochastic gradient descent (SGD)
Gradient computation	$\theta(l)$	$\theta(1)$	$\theta(1)$
Model's updates	$\theta(1)$	$\theta(l)$	$\theta(l)$

Solution (or compromise): batch gradient descent:

- 1) sample a batch instead of a single sample
- 2) compute the gradient on the batch
- 3) update

Result:

- + Gradient computation is still in constant complexity (= batch size). If hardware can parallelize: same time as for 1 element.
- + Number of updates is greatly reduced (divided by batch size w.r.t. SGD).
- + Variance of the estimate of the gradient is reduced using batches.

Optimization and tricks

Stochastic Gradient Descent

In practice, you need to select the batch size.

Small batch size: higher number of updates, high variance of the estimate of the gradient

Large batch size: low number of updates, low variance of the estimate of the gradient

Solution: take very large batch??

No, in practice, gradient's computation too long to compute after a certain threshold.

General solution: take a batch as large as your GPU memory can fit.

Takeaway for Batch SGD:

- Faster than GD
- Faster than SGD with less variance on the gradient estimation
- However, there is still some variance...

Optimization and tricks

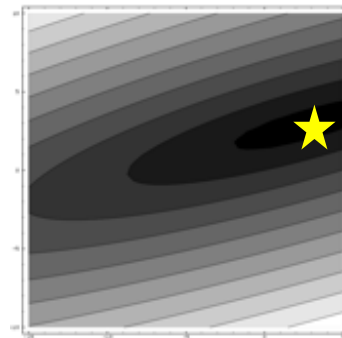

Momentum

Gradient descent (general rule), with ϵ the learning rate:

$$w(t+1) = w(t) - \epsilon \nabla J(f(w(t)))$$

Two problems:

- Learning rate super important!
- Because we only estimate the gradient (when using Batch SGD), it can be noisy (i.e. variance in the estimation).



Optimization and tricks

Momentum

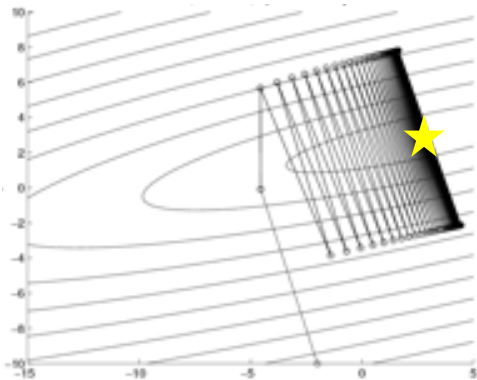
Gradient descent (general rule), with ϵ the learning rate:

$$w(t+1) = w(t) - \epsilon \nabla J(f(w(t)))$$

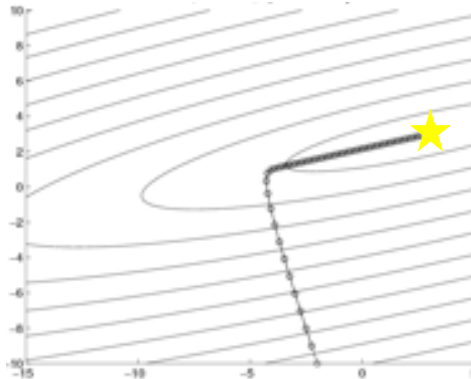
Two problems:

- Learning rate super important!

$\epsilon = 1$: >100 iterations



$\epsilon = 0.1$: 52 iterations



Optimization and tricks

Momentum

Gradient descent (general rule), with ϵ the learning rate:

$$w(t+1) = w(t) - \epsilon \nabla J(f(w(t)))$$

Because we only estimate the gradient (when using Batch SGD), it can be noisy (i.e. variance in the estimation).

Solution: leaky average:

$$\begin{aligned} v(t+1) &= \mu v(t) - \epsilon \nabla J(f(w(t))) \\ w(t+1) &= w(t) + v(t+1) \end{aligned}$$

I.e. Step of the update is an average of the direction given by the gradient, and the previous direction.

To know more:


Sutskever I, Martens J, Dahl G, Hinton G. On the importance of initialization and momentum in deep learning. In International conference on machine learning 2013 Feb 13 (pp. 1139-1147).

Optimization and tricks


Momentum

$$\begin{aligned} v(t+1) &= \mu v(t) - \epsilon \nabla J(f(w(t))) \\ w(t+1) &= w(t) + v(t+1) \end{aligned}$$


$\epsilon = 1$: >100 iterations



$\epsilon = 0.1$: 52 iterations



$\epsilon = 0.1, \mu = 0.5$: 29 iterations



Optimization and tricks

Optimization today

- Using SGD with a fixed learning rate: OK for some problems
- Momentum helps in general
- Out of the scope of this class:
 - learning rate scheduling
 - per-coordinate learning rates
- Adam: in general a good choice. Finding good hyperparameters still a trial & error process.

Kingma, Diederik P., and Jimmy Ba. "Adam: A method for stochastic optimization." arXiv preprint arXiv:1412.6980 (2014).

Optimization and tricks

Questions?




Optimization and tricks


Fighting overfitting

- Neural networks are prone to overfitting


Underfitting



Good fit



Overfitting



Optimization and tricks

Fighting overfitting

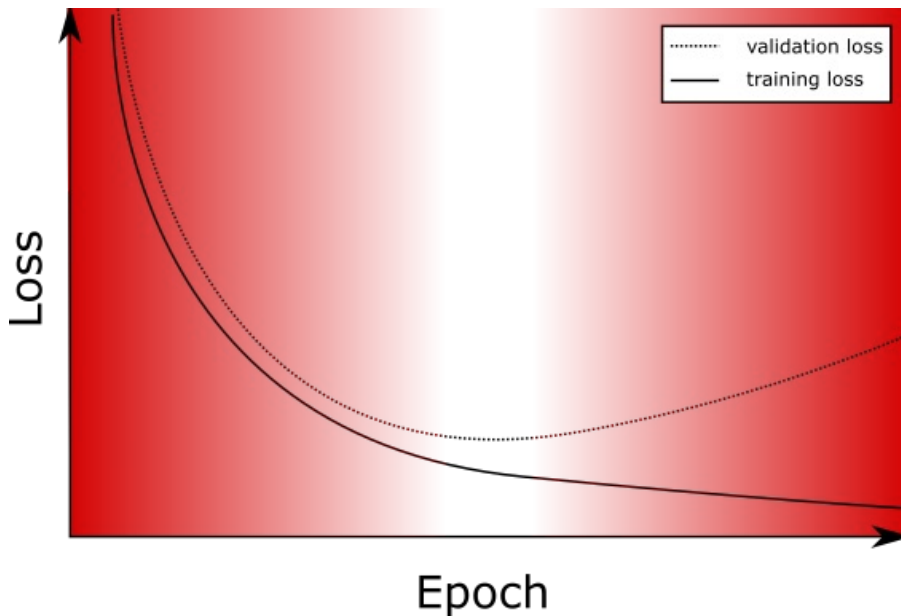
- Neural networks are prone to overfitting
- One reason: they contain a lot of parameters, can model complex functions
- 1st solution: keep the networks small...
- 2nd solution: have more data...
- Other tricks today

Optimization and tricks

Are you overfitting?

Underfitting:

- Train longer
- Augment capacity



Overfitting:

- Stop earlier
- Regularize
- Show more data

Optimization and tricks

Weight decay


- Hypothesis: you have a big capacity w.r.t. number of training samples
- The model can overfit by having weights that adapts to the training samples. In other words, it can model a function that becomes too complex.
- Solution: penalize too complex models!
- Weight decay:
 - Measure the complexity of a model: L2 norm of the weights.
 - Penalize the complexity: add the L2 norm of the weights to the loss
 - Result: weights shrink towards 0

Optimization and tricks


Dropout

- One possible reason from overfitting: co-adaptation
- Several intuitions proposed by authors: role of sex in evolution, spread of conspiracy theories...
- Idea: break co-adaptation by disabling nodes at training time
- For MLP

Srivastava, N., Hinton, G., Krizhevsky, A., Sutskever, I. and Salakhutdinov, R., 2014. Dropout: a simple way to prevent neural networks from overfitting. *The journal of machine learning research*, 15(1), pp.1929-1958.



(a) Standard Neural Net



(b) After applying dropout.

Optimization and tricks

Dropout

- In practice remove a connection with a given probability
- no preferential path can be learned
- slightly different models are learned at every path
- the final model is a kind of average (= ensemble model)

Optimization and tricks

Batchnorm


- Intuition: normalize all your feature maps with the statistics you observe within each batch (in general, before the activation function)
- Normalization is always a good thing
- Having a similar scale for all of the inputs allows the network not to have to deal with that
- It will prevent to have a layer dominating because it has high values -> regularization.
- In practice:
 - At training time: normalize w.r.t. the statistics observed over each batch
 - At test time: normalize w.r.t. the statistics computed over the training set.

Optimization and tricks

Are you overfitting?

Underfitting:

- Train longer
- Augment capacity



Overfitting:

- Stop earlier
- Regularize
- Show more data

Optimization and tricks

Data augmentation

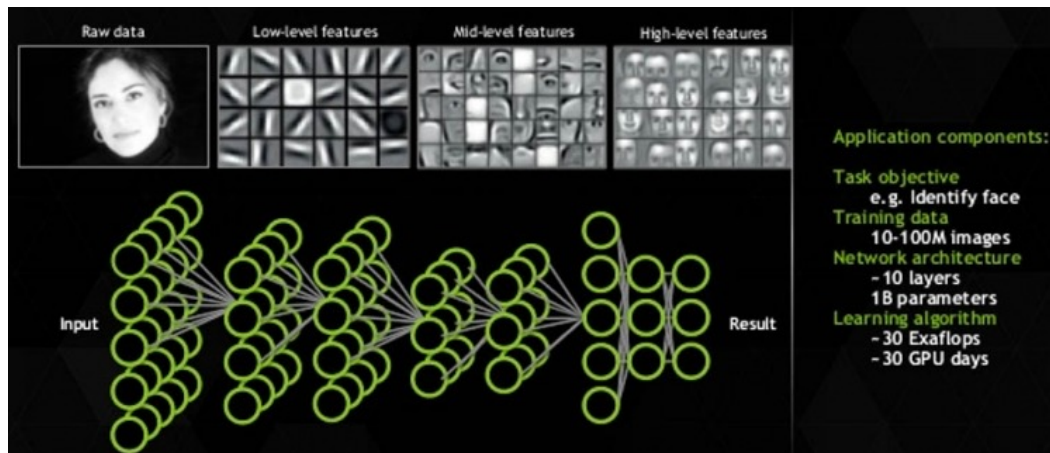
- Set of operations applied randomly to the input data -> obtain different training samples
- Common data augmentation operations for images:
 - Cropping
 - Rotation
 - Vertical/horizontal flips
 - Injecting noise?

Optimization and tricks

Fine tuning

- Common technique to train a model when not enough data
- Intuition: a lot of the filters learnt by a model are applicable in most situations (e.g. edge detectors, texture,...)

Credit: NVIDIA



Optimization and tricks

Fine tuning

- Common technique to train a model when not enough data
- Intuition: a lot of the filters learnt by a model are applicable in most situations (e.g. edge detectors, texture,...)
- Idea: take a model learnt on a big dataset (e.g. ImageNet)
- Replace the layers you need to change (generally last ones), and initialize them
- Re-train (or fine-tune) everything on your data.

Today's menu

- 1 Optimisation
- 2 Sequence modeling
- 3 Visual Question Answering
- 4 Activity

Sequences


What can you do with an image?

Input image



Sequences

What can you do with ~~an image~~ a sequence?



Sequences

What can you do with a sequence?




- Up until now: single object \mathbf{x} as input
- Sequence: collection of T objects: $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_T$
- Tasks:
 - From a sequence, predict a single value y (e.g. crop type, deforestation, sentiment, ...)
 - From a sequence, predict a sequence $y_1, y_2, \dots, y_{T'}$:
 - With 1-to-1 correspondence between \mathbf{x}_1 and y_1 (e.g. presence/absence of forest)
 - Without 1-to-1 correspondence (e.g. machine translation)

Sequences

What can you do with a sequence?

- Examples of sequences:

- Text
- Stock market
- Sequence of images
- DNA...



Sequences

Model a sequence

- Model the sequence = model $P(x_1, x_2, \dots, x_T)$
- We have:

$$P(x_1, x_2, \dots, x_T) = \prod_{t=1}^T P(x_t \mid x_1, \dots, x_{t-1}).$$

- If we add a Markovian property:
 - order 0: $P(x_1, x_2, x_3) = P(x_1)P(x_2)P(x_3)$
 - order 1: $P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_2)$
 - order 2: $P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_2, x_1)$
 - ...

Sequences

Model a sequence

- order 1: $P(x_1, x_2, x_3) = P(x_1)P(x_2|x_1)P(x_3|x_2)$
- bigram can be estimated from a corpus:
 - $P(x_2|x_1) = \frac{n(x_2, x_1)}{n(x_1)}$
- Trade-off between modeling long-term dependencies and frequency of co-occurrences

Sequences

Model a sequence

- Ideally, we want to keep the whole sequence in the modeling of each time step, i.e.

$$P(x_1, x_2, \dots, x_T) = \prod_{t=1}^T P(x_t \mid x_1, \dots, x_{t-1}).$$

- Computationally impossible to model: too many parameters...

Sequences

Recurrent Neural Network

- Main idea of RNN:

$$P(x_t \mid x_{t-1}, \dots, x_1) \approx P(x_t \mid h_{t-1}),$$

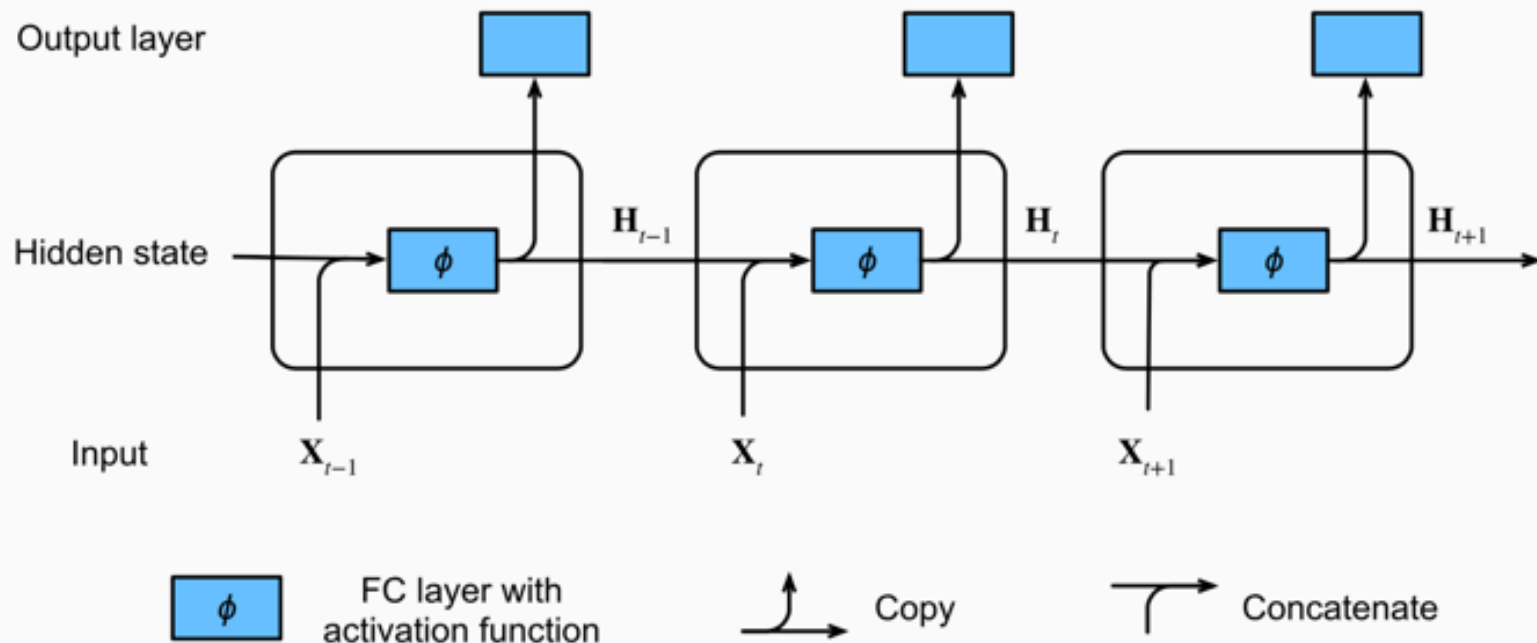
- with:

$$h_t = f(x_t, h_{t-1}).$$

- The variable h_{t-1} is a hidden state: it stores information on the sequence until step $t - 1$

Sequences

Recurrent Neural Network



Sequences

Dealing with text

- A RNN expect a vector of numbers as an input.
- A text is a string
- Tokenization:
 - “Hello everybody, enjoy the class and enjoy life” ->
 - [“Hello”, “everybody”, “enjoy”, “the”, “class”, “and”, “enjoy”, “life”]
- Words converted to numbers using a vocabulary:
 - Either created from the text: [0, 1, 2, 3, 4, 5, 2, 6]
 - pre-defined (from a large corpus, to use pre-trained models)
- Reversible operation

Sequences

The problem with this RNN

- Hard to learn long-term dependencies because of exploding and vanishing gradient
- Desirable properties:
 - Important information should be retained after a while (long-term memory)
 - All the information from the recent past should be stored (short-term memory)
- Intuition: the hidden state cannot be computed as a uniform average of all the previous steps

Sequences

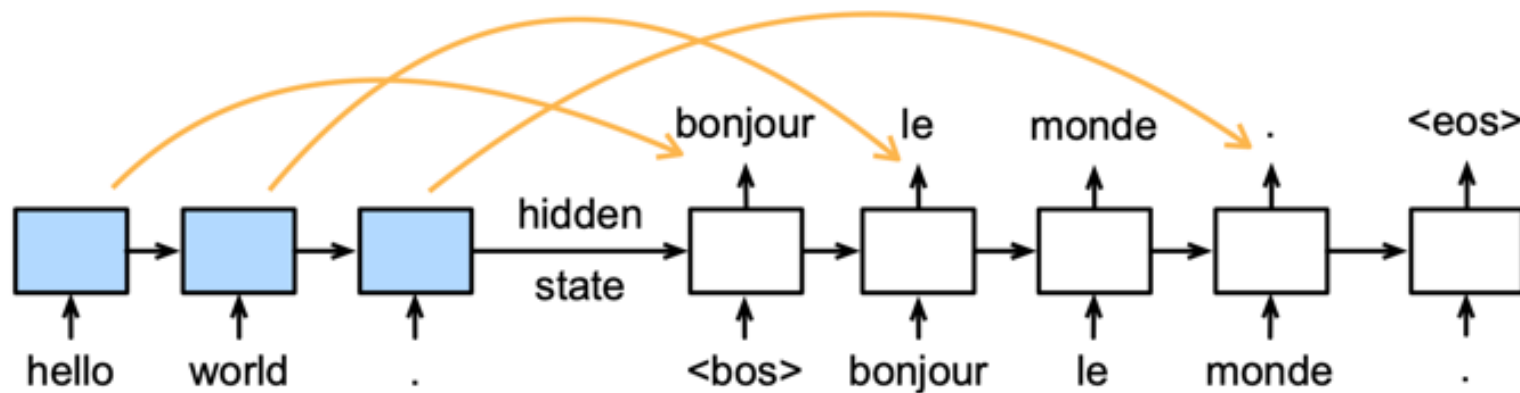
Gates

- LSTM (Long-short-term memory) introduces the concept of gate:
- A gate = learnable parameter deciding how much information we want to keep
- Regrouped in a memory cell in charge of computing the gates from an internal state
- In LSTM, 3 gates:
 - input gate: how much the input (x_t) should influence the current step?
 - forget gate: how much of the internal state should we keep?
 - output gate: how much the current step should influence the output?

Sequences

Attention

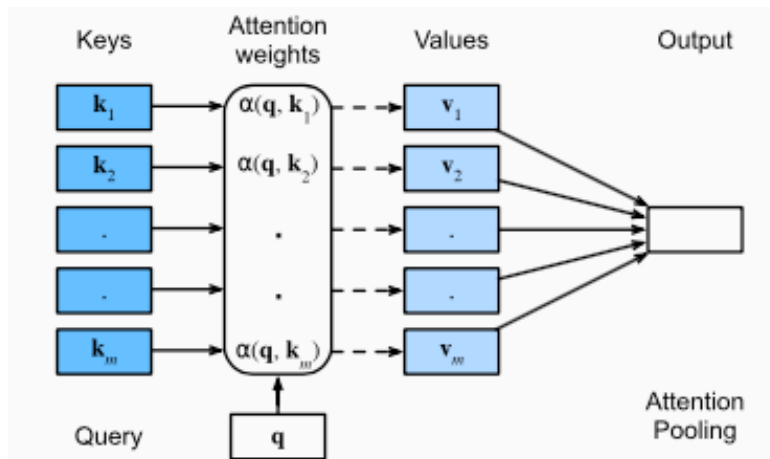
- Let's take an example of sequence to sequence translation.
- Generated tokens are related to a combination of input tokens



Sequences

Attention

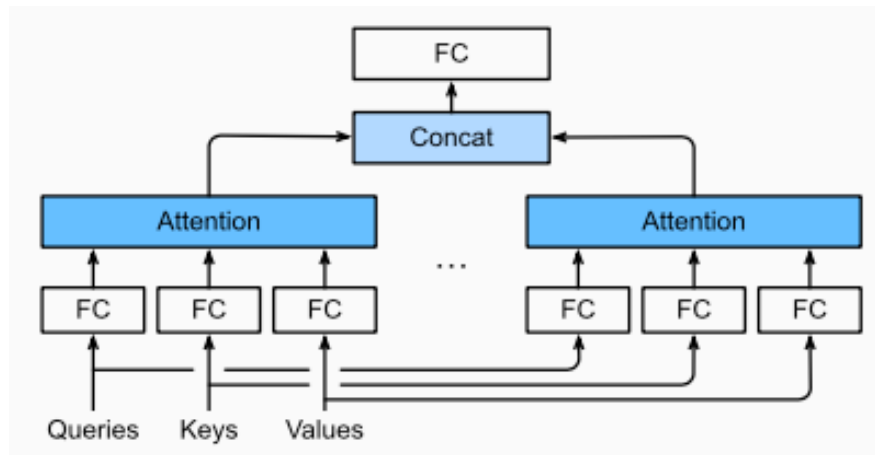
- Attention = make a linear combination of values, based on the combability between a set of keys and a query
- Allows to select relevant input tokens to generate output tokens



Sequences

Attention


- In practice, it can be useful to capture several dependencies between tokens (e.g. long vs short range).
- This can be done by allowing different representations of queries, keys and values



Sequences




Self-attention

- Let's feed a sequence of tokens to a model
- Let's give each token its own set of query, keys and values -> each token can attend (through its query) any combination of others tokens (through their keys).



Sequences


Self-attention



Sequences

Transformer

- Self-attention + positional encoding -> Transformer architecture



Sequences

Conclusion

- Very basics of sequence modeling
- RNN still used today, outperformed by transformers
- Transformers allow for interactions of each token with any other -> computationally heavy, but much better modeling of interactions
- Transformers can also be used for vision

Today's menu

- 1 Optimisation
- 2 Sequence modeling
- 3 Visual Question Answering
- 4 Activity


Visual Question Answering

Visual Question Answering

- Objective: provide an answer (in natural language) from:
 - One image
 - One question (in natural language)



What is on the coffee table ?
candles



What color is the hydrant ?
black and yellow



What is on the bed ?
books



What is the long stick for ?
whipping

Samples from the VQA2 dataset (from Teney, 2017))

Visual Question Answering

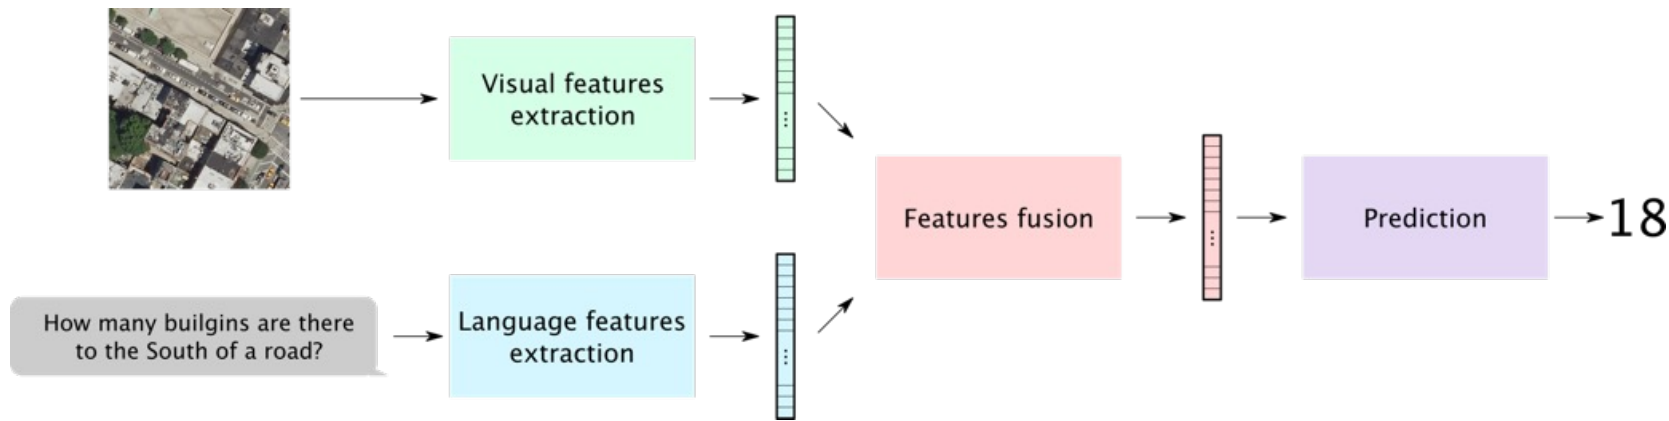
Visual Question Answering

- This is a new task! Introduced by [1] in 2015
- Can be seen as an extension of the Turing test to the visual domain
- Questions are not limited: can cover most of the computer vision tasks
- Can make computer vision results accessible to new people
- Limited applications today. Most studied one: answering question from visually impaired people

Visual Question Answering

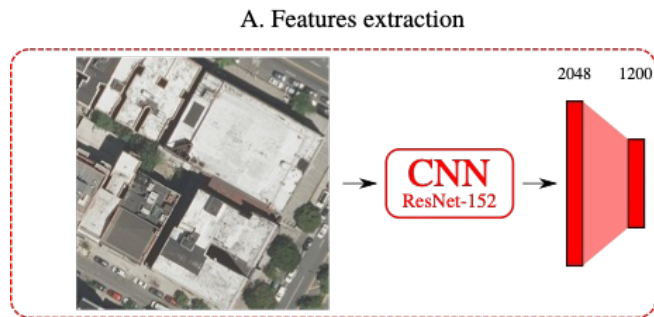
Model

- Models look like this



Visual Question Answering

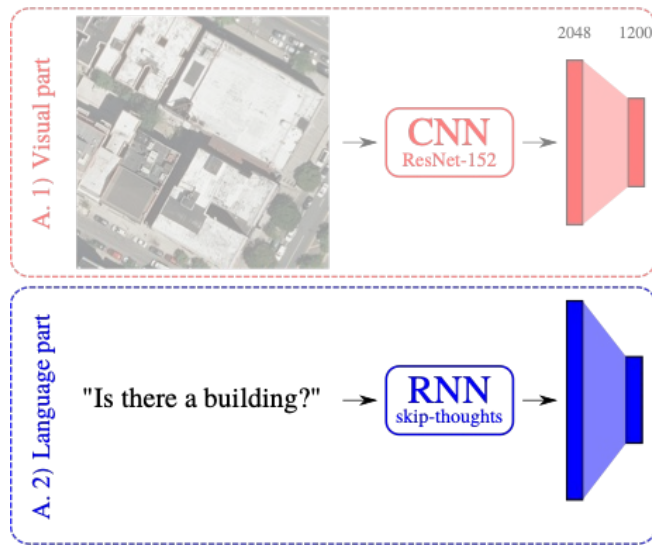
Model



- ResNet-152 pre-trained on ImageNet with a fully-connected layer.

Visual Question Answering

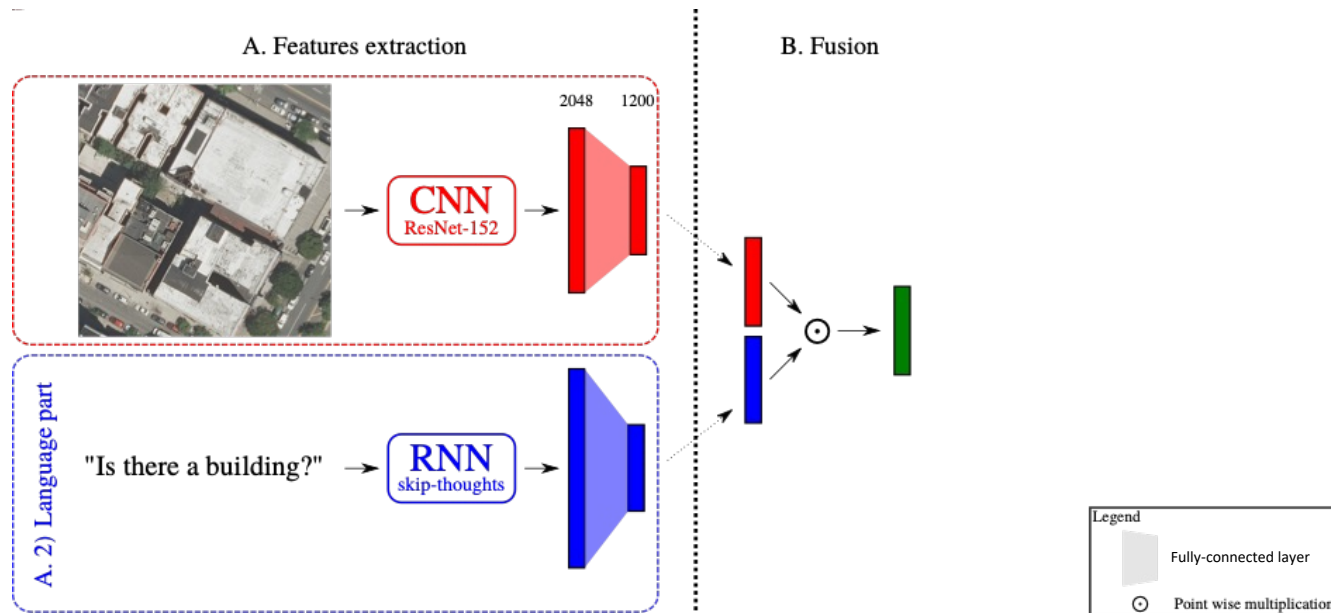
Model



- Sequence encoder based on skip-thoughts : predict the previous and following sentences of a sentence in a book. Trained on BookCorpus.

Visual Question Answering

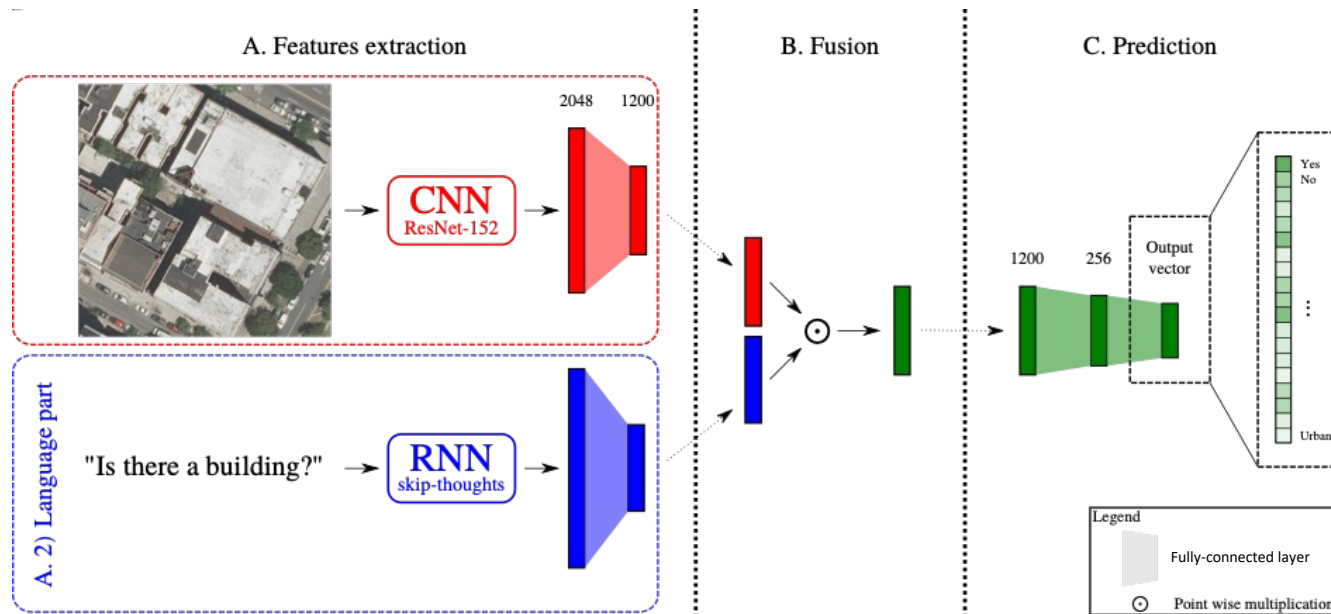
Model



- Vectors fusion by point-wise multiplication.

Visual Question Answering

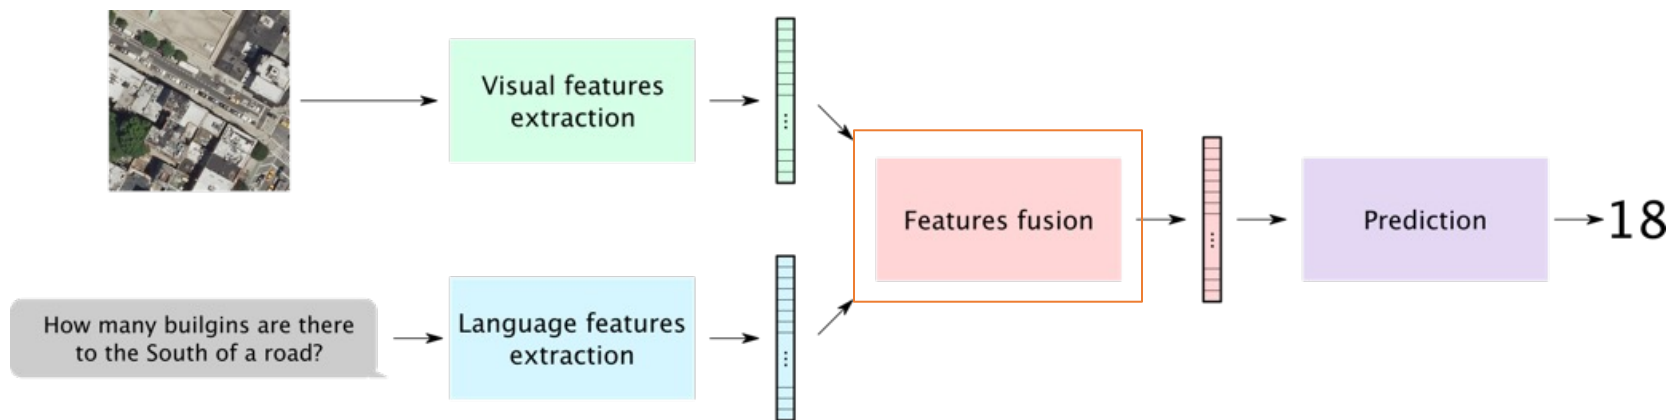
Model



- Predict the most probable answer among a set of pre-defined ones.

Visual Question Answering

Model



Visual Question Answering

Fusion

- Ideally we want to multiply both vectors
- Computationally intractable
- Possibility to use random projections [1], tensor decomposition [2]

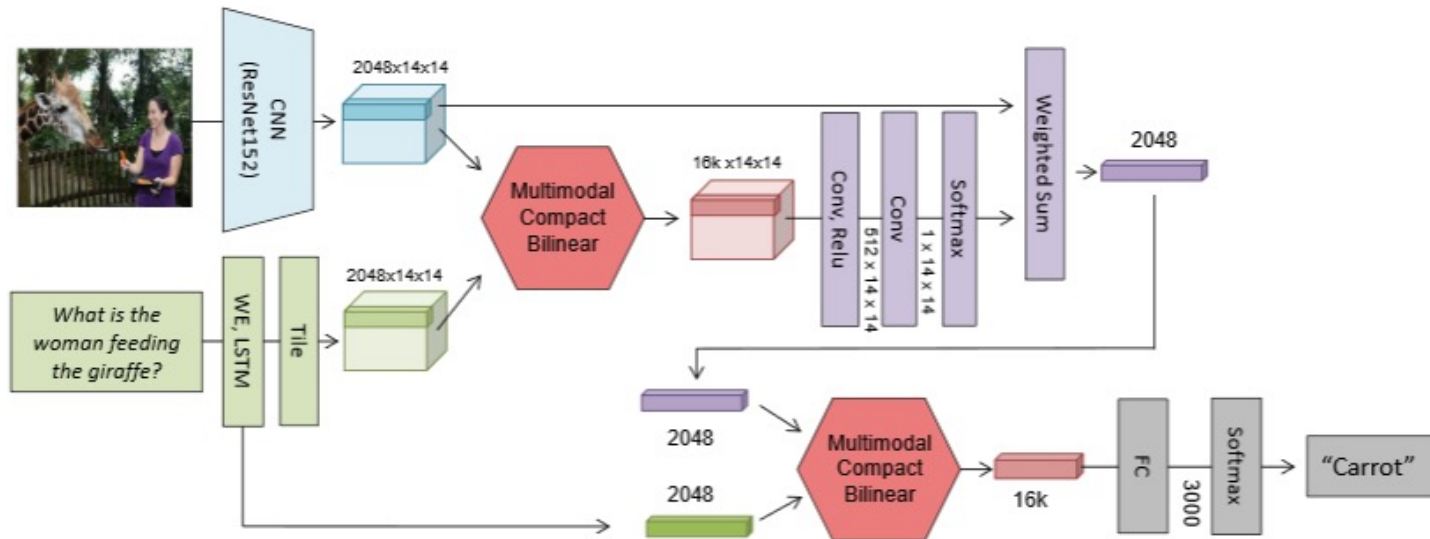
[1] Fukui, Akira, et al. "Multimodal compact bilinear pooling for visual question answering and visual grounding." arXiv preprint arXiv:1606.01847 (2016).

[2] Ben-Younes, Hedi, et al. "Mutan: Multimodal tucker fusion for visual question answering." Proceedings of the IEEE international conference on computer vision. 2017.

Visual Question Answering

Fusion – MCB model

- Fusion at two level: attention and prediction



A first simple solution

Towards new application

RSVQA LR

772 images (Sentinel 2)

77'232 questions



Example:

What is the number of water areas? 7

RSVQA HR

10'659 images (USGS, 15cm)

955'664 questions



Example:

What is the amount of buildings? 7

RSVQAxBEN

590'326 images (Sentinel 2)

14'758'150 questions



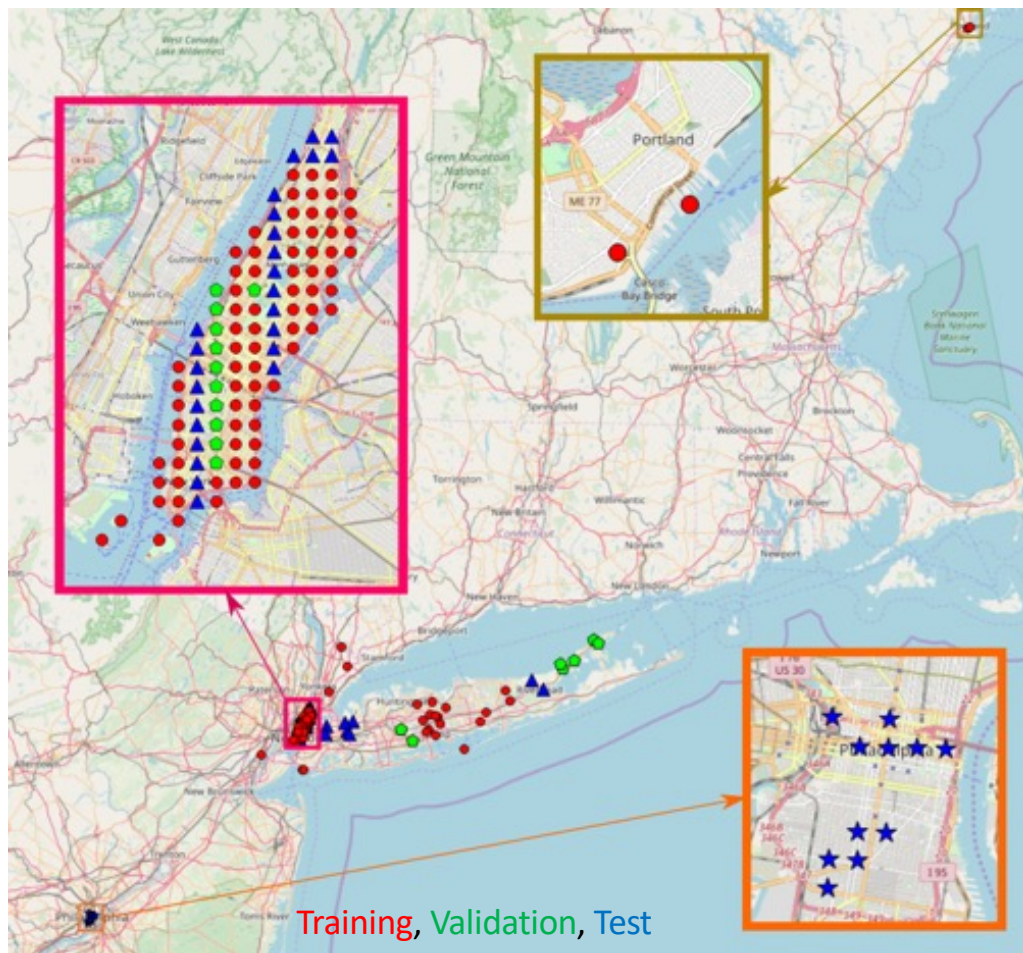
Example:

Are there artificial areas and agricultural areas or water bodies? Yes

A first simple solution

RSVQA HR

- East cost of the US
- 10'659 images from 161 orthophotos
- Two test sets to test spatial generalization.



A first simple solution

Results


- Difficult to count
- Loss in performances on a new area

Type	Accuracy Test set 1	Accuracy Test set 2
Count	68.63% (0.11%)	61.47% (0.08%)
Presence	90.43% (0.04%)	86.26% (0.47%)
Comparison	88.19% (0.08%)	85.94% (0.12%)
Area	85.24% (0.05%)	76.33% (0.50%)
AA	83.12% (0.03%)	77.50% (0.29%)
OA	83.23% (0.02%)	78.23% (0.25%)

A first simple solution

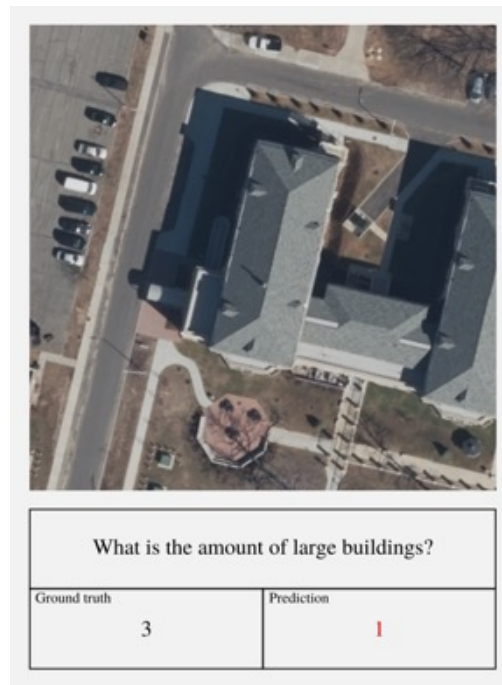
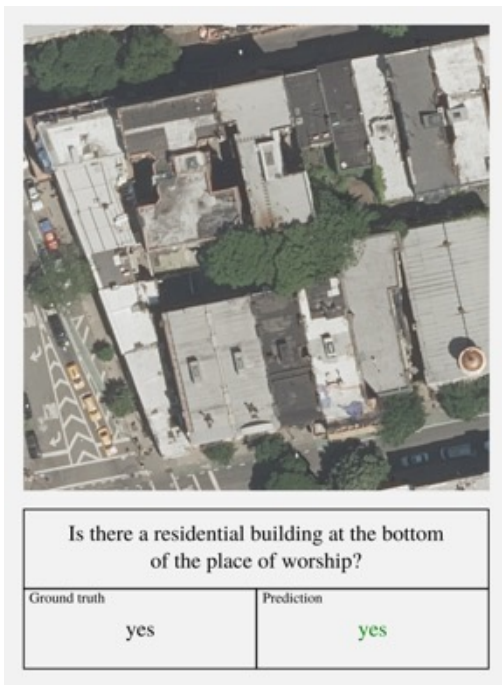
Results

- Difficult to count
- Loss in performances on a new area
- When the answer is wrong, it is still logical



A first simple solution

Visual results



Visual Question Answering

Conclusion

- VQA is a new task
- Can open new usages of computer vision
- A very active research community on the topic